

Computational efficiency: inference latency and memory footprint

