

Inference latency (CPU, single prediction)

